

# Linux内核编程03期：系统调用

文档配套视频地址：<https://wanglitao.taobao.com>  
专注嵌入式精品教程：[www.zhaixue.cc](http://www.zhaixue.cc)

kill() 系统调用

系统调用号

```
arch/arm/include/generated/uapi/asm/unistd-common.h

#define __NR_kill  (__NR_SYSCALL_BASE + 37)
```

系统调用函数实现

```
kernel/signal.c :

SYSCALL_DEFINE2(kill, pid_t, pid, int, sig)
{
    struct kernel_siginfo info;
    prepare_kill_siginfo(sig, &info);
    return kill_something_info(sig, &info, pid);
}

展开后相当于：
asmlinkage long sys_kill(pid_t pid, int sig)
```

系统调用函数声明

```
include/linux/syscalls.h

asmlinkage:GCC扩展，表示读取的参数来自栈中，而非寄存器
/* kernel/signal.c */
asmlinkage long sys_restart_syscall(void);
asmlinkage long sys_kill(pid_t pid, int sig);
asmlinkage long sys_tkill(pid_t pid, int sig);
```

```
arch/arm/kernel/entry-common.S : 保护现场，获取系统调用号
ENTRY(vector_swi)
addne scno, r7, #__NR_SYSCALL_BASE @ put OS number in
ldr tbl, sys_call_table
...
invoke_syscall tbl, scno, r10, __ret_fast_syscall
add r1, sp, #S_OFF
2: cmp scno, #(__ARM_NR_BASE - __NR_SYSCALL_BASE)
eor r0, scno, #__NR_SYSCALL_BASE @ put OS number back
bcs arm_syscall
mov why, #0 @ no longer a real syscall
b sys_ni_syscall @ not private func
...
9001:
sub lr, saved_pc, #4
str lr, [sp, #S_PC]
get_thread_info tsk
b ret_fast_syscall 回到用户态kill，继续执行用户态代码
ENDPROC(vector_swi)

syscall_table_start sys_call_table
#define COMPAT(nr, native, compat) syscall nr, native
#ifdef CONFIG_AEABI
#include <calls-eabi.S>
#else
#include <calls-oabi.S>
#endif
#undef COMPAT
syscall_table_end sys_call_table

#define NATIVE(nr, func) syscall nr, func
```

```
arch/arm/include/generated/calls-eabi.S :

NATIVE(0, sys_restart_syscall)
NATIVE(1, sys_exit)
NATIVE(2, sys_fork)
NATIVE(3, sys_read)
NATIVE(4, sys_write)
NATIVE(5, sys_open)
NATIVE(6, sys_close)
NATIVE(8, sys_creat)
NATIVE(9, sys_link)
NATIVE(10, sys_unlink)
NATIVE(11, sys_execve)
NATIVE(12, sys_chdir)
NATIVE(14, sys_mknod)
NATIVE(15, sys_chmod)
NATIVE(16, sys_lchown16)
NATIVE(19, sys_lseek)
NATIVE(20, sys_getpid)
NATIVE(21, sys_mount)
NATIVE(23, sys_setuid16)
NATIVE(24, sys_getuid16)
NATIVE(26, sys_ptrace)
NATIVE(29, sys_pause)
NATIVE(33, sys_access)
NATIVE(34, sys_nice)
NATIVE(36, sys_sync)
NATIVE(37, sys_kill)
NATIVE(38, sys_rename)
NATIVE(39, sys_mkdir)
其实就是定义一个函数入口指针 .long sys_kill
```

```
tools/include/nolibc/nolibc.h:
#define my_syscall2(num, arg1, arg2) \
({ \
    __register long _num asm("r7") = (num); \
    __register long _arg1 asm("r0") = (long)(arg1); \
    __register long _arg2 asm("r1") = (long)(arg2); \
    __asm volatile ( \
        __asm "svc #0\n" \
        __asm "r"(_arg1), "r"(_arg2), \
        __asm "r"(_num) \
        __asm "memory", "cc", "lr" \
    ); \
    __arg1; \
})

static __attribute__((unused))
int sys_kill(pid_t pid, int signal)
{
    __return my_syscall2(__NR_kill, pid, signal);
}

static __attribute__((unused))
int kill(pid_t pid, int signal)
{
    __int ret = sys_kill(pid, signal);

    __if (ret < 0) {
        __SET_ERRNO(-ret);
        __ret = -1;
    }
    __return ret;
}
```

```
/usr/arm-linux-gnueabi/lib/libc-2.31.so:

000dad70 <syscall@@GLIBC_2.4>:
dad70:  e1a0c00d  __mov__ip, sp
dad74:  e92d00f0  __push__ {r4, r5, r6, r7}
dad78:  e1a07000  __mov__r7, r0
dad7c:  e1a00001  __mov__r0, r1
dad80:  e1a01002  __mov__r1, r2
dad84:  e1a02003  __mov__r2, r3
dad88:  e89c0078  __ldm__ip, {r3, r4, r5, r6}
dad8c:  ef000000  __svc__0x00000000
dad90:  e8bd00f0  __pop__ {r4, r5, r6, r7}
dad94:  e3700a01  __cmn__r0, #4096; 0x1000
dad98:  312fff1e  __bxcc__lr
dad9c:  eafcf2c3  __b__178b0 <__libc_start_main@@GLIBC_2.4+0x278>
```

```
/usr/arm-linux-gnueabi/lib/libc.a:

00000000 <_kill>:
0:  e52d7004  __push__ {r7}; (str r7, [sp, #-4]!)
4:  e3a07025  __mov__r7, #37; 0x25
8:  ef000000  __svc__0x00000000
c:  e49d7004  __pop__ {r7}; (ldr r7, [sp], #4)
10: e3700a01  __cmn__r0, #4096; 0x1000
14: 312fff1e  __bxcc__lr
18:  eaffffff  __b__0 <__syscall_error>
```

SVC指令，即以前的SWI指令，软中断